

Appendix B

Algorithm 1: Duplicate and Near-Duplicate Image Detection

Input: image_list, similarity_thresh

Output: duplicate_pairs, near_duplicate_pairs, num_duplicates, num_near_duplicates

Main Detection Procedure

1. Initialize empty lists for duplicate_pairs and near_duplicate_pairs
2. For each pair (i, j) where $0 \leq i < j < \text{length}(\text{image_list})$:
 - If $\text{image_list}[i]$ and $\text{image_list}[j]$ are exactly identical:
 - Add (i, j) to duplicate_pairs
 - Else:
 - similarity_score = ORB_SIMILARITY($\text{image_list}[i]$, $\text{image_list}[j]$)
 - If $\text{similarity_score} > \text{similarity_thresh}$ and < 1.0 :
 - Add (i, j) to near_duplicate_pairs
3. Count duplicates using grouping:
 - num_duplicates = COUNT_DUPLICATES(duplicate_pairs)
 - num_near_duplicates = COUNT_DUPLICATES(near_duplicate_pairs)
4. Return all pairs and counts

Function: ORB_SIMILARITY(img1, img2)

1. Detect ORB keypoints and descriptors for both images
2. Match descriptors using brute-force matcher
3. If no matches found: return 0
4. Count strong matches ($\text{distance} < 50$)
5. Return ratio: $\text{strong_matches} / \text{total_matches}$

Function: COUNT_DUPLICATES(pairs)

1. Initialize a dictionary groups, creating a key for every unique element found in all pairs.
2. For each key in groups and each pair (i, j) :
 - If key matches i or j : add both to groups[key]
3. For each group: keep all except smallest after sorting

4. Return count of unique extracted elements